

# A Language's Unigram Entropy Distribution Predicts Self-Paced Reading Times

Josef Klafka

jklafka@uchicago.edu  
Department of Psychology  
University of Chicago

Daniel Yurovsky

yurovsky@uchicago.edu  
Department of Psychology  
University of Chicago

## Abstract

We provide evidence against the Uniform Information Density (Jaeger, 2010) and wrap-up effect hypotheses.

**Keywords:** Entropy; self-paced reading; information theory; language processing

## Uniform Information Density

How do people convey information in speech and text? Is there a regular distribution to how information is conveyed across speech and text? Does speech, with its phonological properties, different from writing, with its semiotic properties, in how information is conveyed?

A major line of research in the current century has asked how speakers and writers distribute information in the linguistic material they produce. This begins with Claude Shannon in the years after World War II. Shannon (1948) defined “information” as a “reduction in uncertainty”. Shannon followed by quantifying uncertainty in his measure of *entropy*: the amount of uncertainty on the outcome of a random variable. Shannon proposed that the most efficient method of sending information through a noisy channel is at a constant rate. Genzel & Charniak (2002) apply this principle to human communication: if people communicate information through speech and text optimally, then we should transmit information to one another at a constant rate. Genzel & Charniak created a distribution of estimates for sentence-level entropy, each sentence taken without context, for newspaper articles from the Penn Treebank corpus. Obtaining a roughly monotonically increasing linear distribution, Genzel & Charniak proposed the *Constant Entropy Rate* (CER) principle: the entropy rate in speech and text with context should be constant as sentence number increases. Aylett & Turk (2004) proposed a similar principle for the phonological level.

Jaeger (2010) extends the CER principle to all levels of human communication: the *Uniform Information Density* (UID) hypothesis. Jaeger argues that speakers will try to distribute information more evenly over the course of an utterance, to be as close to channel capacity (in a Shannon sense) as possible. In particular, Jaeger singles out the production of “that” at the beginning of relative clauses in English and the use of contractions such as “he’s” and “you’re” versus “he is” and “you are” as instances where higher information density in the clause elicits the production of more material by speakers.

The UID hypothesis

## UID challenges

More recent work has stood out in contrast to the traditional UID perspective. Zhan and Levy (2018) study Mandarin Chi-

nese classifier use, and find that the use of specific classifier, versus general classifiers, appears more often in cases where the production of the corresponding noun is more difficult than when the production is easier, as would be predicted by UID. The UID perspective is challenged in Yu et al. (2016), by performing an analysis of entropy by position in the text portion of the British National Corpus. They use the following formula for each word position  $X$  of sentences of fixed length  $k$  from the corpus, where each  $i$  is a word occurring in position  $X$  and  $p_i$  is the number of times word  $i$  occurs in position  $X$  divided by the number of total words that occur in position  $X$  i.e. the number of sentences of length  $k$ .

$$H(X) = \sum_w p(w) \log(p(w))$$

Yu et al. (2016) refer to their distribution for English as a ‘three-step distribution’: relatively low entropy at the beginning of a sentence, then a jump, then flat entropy in the middle, a dip before the final position and a jump with the final word. View the figure below for a visual demonstration.

## Methods

We used the CHILDES TalkBank (Brown 1973; MacWhinney 2000) corpora database of spoken adult-child conversations. We first used the Brown and Providence English corpora from CHILDES. The Brown corpus contains conversations between three young children (over 1.5 years old and under 6 years old) and their families in the home. The Providence corpus contains transcriptions of audio/video files which recorded interactions between children between 1 and 3 years old and their parents in the home. We divided each corpus by speaker into child and non-child categories. We further divided the corpora by utterance length, so that all sentences of length  $k$  (e.g. 6) were grouped together. Finally, within each utterance length, we computed the unigram entropy measure for each position.

The unigram entropy measure works as

Figure for the method?

Plots for the entropy distributions are below. The adult and child unigram entropy distributions track one another extremely well, with the child distribution for each sentence length having relatively lower entropy.

We also ran this analysis on Spanish, German, Mandarin and Japanese corpora from CHILDES. The Spanish corpus we used was the XXX and the German corpus we used was the XXX. The Mandarin corpus we used was the XXX and the Japanese corpus we used was the XXX. For Mandarin, we used pinyin transliterations of the utterances in the corpus

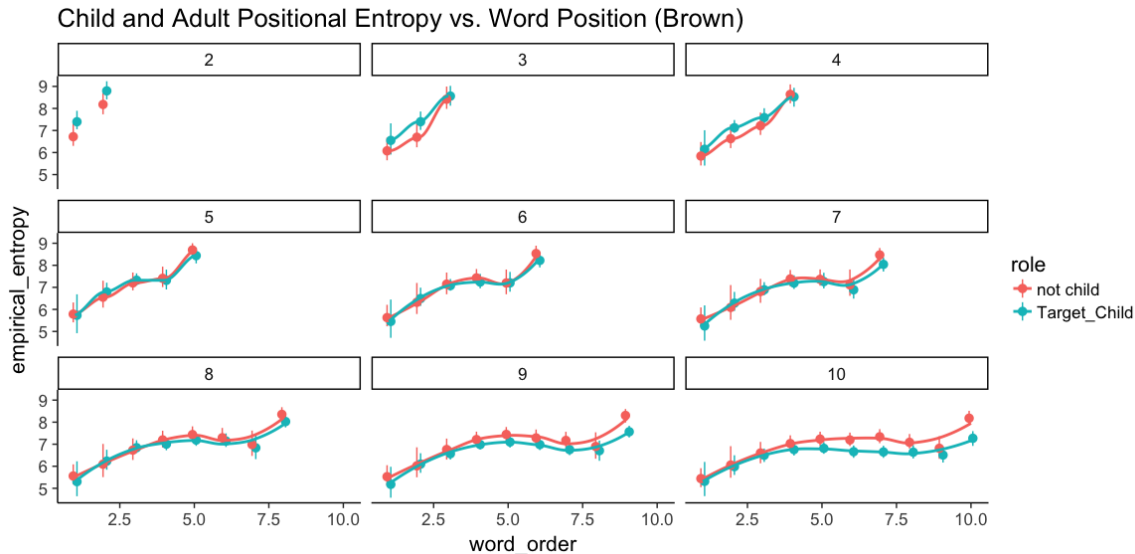


Figure 1: Brown corpus entropy

with demarcated word boundaries, and for Japanese we used romanji transliterations of words in the corpus. Japanese Hiragana, Katakana and Kanji writing systems, and the Chinese characters used for writing Mandarin do not normally demarcate word boundaries by spacing words apart, and for normal Japanese and Chinese writing including spaces between word boundaries can have a negative effect on reading times (Sainio et al, 2007; Bai et al, 2008).

## Results

We found a distinct three-step distribution for English, Spanish and German CHILDES corpora, with a slight dip in the penultimate position of each sentence. The Mandarin and Japanese data, by comparison,

## Analysis

### Wikipedia Data

We also wanted to apply this analysis to large-scale text data, to compare with the parent-child speech data results from using CHILDES. Using Giuseppe Attardi’s Wikiextractor tool <sup>1</sup>, we extracted the text corpora for 165 languages from Wikipedia by downloading the data dump of Wikipedia. Each language corpus was cleaned and limited to sentences between 6 and 50 words. Similar to the data from CHILDES, we divided each corpus by sentence length, and then computed the unigram entropy measure on each word position

within each sentence length. HOW DID WE JOIN TOGETHER SENTENCE LENGTHS

## Clustering

### Conclustions and future research

The *wrap-up effect* is a popular hypothesis within the eye-tracking community. The hypothesis states that, in written text, sentence-final words are processed more slowly on average than sentence-medial or sentence-initial words, due to readers integrating information from the entire sentence in order to form a final, coherent thought expressed by the sentence, among other reasons. With these two ideas, information production distribution in spoken and written communication, and information integration in reading written texts, we move on to discuss cross-linguistic differences in the distribution of information across sentences.

## Acknowledgements

Place acknowledgments (including funding information) in a section at the end of the paper.

## References

<sup>1</sup><https://github.com/attardi/wikiextractor>




Figure 2: Providence corpus unigram entropy